

## ▼ Data Reading and Preprocessing

```
1 import os
2 import copy
3 import itertools
4 import numpy as np
5 import seaborn as sns
6 import scipy
7 import tables as tb
8 from mpl_toolkits.mplot3d import Axes3D
9 import warnings
10 warnings.filterwarnings('ignore')
11
12 %matplotlib inline
13 %pylab inline
```

📄 Populating the interactive namespace from numpy and matplotlib

```
1 # Reading data
2 data = os.path.join('drive', 'My Drive', 'dl_data', 'QIS_EXAM_200Events')
3 test_input = np.load(os.path.join(data, 'test_input.npy'), allow_pickle=True)
4 train_input = np.load(os.path.join(data, 'training_input.npy'), allow_pickle=True)
5
6 test_input = test_input[()]
7 train_input = train_input[()]
8
9 # Segregating the test and train data
10 y_train = np.array([np.array([0])*50 + [np.array([1])*50])
11 y_test = np.array([np.array([0])*50 + [np.array([1])*50])
12
13 x_train = np.concatenate([train_input['0'], train_input['1']])
14 x_test = np.concatenate([test_input['0'], test_input['1']])
```

```
1 # Shuffling the dataset
2
3 def get_indices(_seed, _len):
4     np.random.seed(_seed)
5     indices = np.random.permutation(_len)
```

```

5         indices = np.random.permutation(_len_)
6         return indices
7
8     indices_train = get_indices(0,len(x_train))
9     indices_test = get_indices(1,len(x_test))
10
11     x_train = x_train[indices_train]
12     y_train = y_train[indices_train]
13
14     x_test = x_train[indices_test]
15     y_test = y_test[indices_test]
16
17     # Standardize features by removing the mean and scaling to unit variance
18     scaler = StandardScaler()
19     x_train_scale = scaler.fit_transform(x_train)
20     x_test_scale = scaler.fit_transform(x_test)

```

## ▼ Exploratory Data Analysis

```

1     # 3D visualization of data points
2
3     CMAP = sns.diverging_palette(220, 20, s=99, as_cmap=True, n=2500)
4
5     def plot3D(X, target, elev=0, azimuth=0, title=None, sub=111):
6         x = X[:, 0]
7         y = X[:, 1]
8         z = X[:, 2]
9
10        fig = plt.figure(figsize=(12, 8))
11        ax = Axes3D(fig)
12        mappab = ax.scatter(x, y, z, c=target, cmap=CMAP)
13
14        if title is not None:
15            ax.set_title(title)
16        ax.set_xlabel('Component 1')
17        ax.set_ylabel('Component 2')
18        ax.set_zlabel('Component 3')
19
20        # to change your point of view
21        ax.view_init(elev=elev, azimuth=azimuth)
22        fig.colorbar(mappable=mappab, label='Target variable')


```

```
22 fig.colorbar(mappable=mappab, label= 'target variable' )
23 plt.show()
24
```


```
1 # Visualizing all the combinations of data to get a clear picture of classes
2
3 y = y_train.squeeze(1)
4 for i in range(5):
5     for j in range(i+1,5):
6         for k in range(j+1,5):
7             arr = x_train[:,[i,j,k]]
8             plot3D(arr, y, elev=15, azim=50, title=str([i,j,k]))
```




[0, 1, 2]





[0, 1, 3]










[0, 1, 4]







[0, 2, 3]











[1, 2, 3]










[1, 2, 4]












```


1  # signal basetrack distribution along the axis
2  axis = 'X'
3
4  fig = plt.figure(figsize = [20, 10])
5  fig.add_subplot(221)
6  plt.hist(x_train[y==1][:,[0]], bins=500, histtype='step')
7  fig.add_subplot(222)
8  plt.hist(x_train[y==1][:,[1]], bins=500, histtype='step')
9  fig.add_subplot(223)
10 plt.hist(x_train[y==1][:,[2]], bins=500, histtype='step')
11 fig.add_subplot(224)
12 plt.hist(x_train[y==1][:,[3]], bins=500, histtype='step')
13 plt.show()
14 fig.add_subplot(321)
15 fig = plt.figure(figsize = [10, 5])

```

```
16 plt.hist(x_train[y==1][:,[4]], bins=500, histtype='step')
17 plt.show()
```













```
1 # background basetrack distribution along the axis
2 axis = 'X'
3
4 fig = plt.figure(figsize = [20, 10])
5 fig.add_subplot(221)
6 plt.hist(x_train[y==0][:,[0]], bins=500, histtype='step')
7 fig.add_subplot(222)
8 plt.hist(x_train[y==0][:,[1]], bins=500, histtype='step')
9 fig.add_subplot(223)
10 plt.hist(x_train[y==0][:,[2]], bins=500, histtype='step')
11 fig.add_subplot(224)
12 plt.hist(x_train[y==0][:,[3]], bins=500, histtype='step')
13 plt.show()
14 fig.add_subplot(321)
15 fig = plt.figure(figsize = [10, 5])
16 plt.hist(x_train[y==0][:,[4]], bins=500, histtype='step')
17 plt.show()
```







## ▼ Classification

```
1 import xgboost as xg
2 from xgboost import XGBClassifier
3 from sklearn.model_selection import StratifiedKFold, GridSearchCV
4 from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier
5 from sklearn.tree import DecisionTreeClassifier
6 from sklearn.preprocessing import StandardScaler, Normalizer
7 from sklearn import svm
8 from sklearn import metrics
9 from sklearn.metrics import (
10     classification_report,
11     confusion_matrix, roc_curve,
12     roc_auc_score,
13     log_loss
14 )
15
16 from keras.layers.core import Dense, Activation, Dropout
17 from keras.models import Sequential
18 from keras.optimizers import Adam
19 from keras.utils import np_utils
20 from keras.callbacks import EarlyStopping, ModelCheckpoint
21
```

```
1 # utility function for plotting ROC-AUC
2 def plot_metrics(y_true, y_pred):
3     fpr, tpr, thresholds = roc_curve(y_true, y_pred)
4     roc_auc = roc_auc_score(y_true, y_pred)
5
6     plt.plot(fpr, tpr, label='ROC AUC=%f' % roc_auc)
7     plt.xlabel("FPR")
8     plt.ylabel("TPR")
9
```



```

9     plt.legend()
10     plt.title("ROC Curve")
11
12     # Utility function to add noise
13     def add_noise(array, level=0.15, random_seed=34):
14         numpy.random.seed(random_seed)
15         return level * numpy.random.random(size=array.size) + (1 - level) * array
16
17     # Utility function to plot the confusion matrix
18     def plot_confusion_matrix(cm, classes,
19                             normalize=False,
20                             title='Confusion matrix',
21                             cmap=plt.cm.Blues):
22         """
23         This function prints and plots the confusion matrix.
24         Normalization can be applied by setting `normalize=True`.
25         """
26         plt.imshow(cm, interpolation='nearest', cmap=cmap)
27         plt.title(title)
28         plt.colorbar()
29         tick_marks = np.arange(len(classes))
30         plt.xticks(tick_marks, classes, rotation=45)
31         plt.yticks(tick_marks, classes)
32
33         if normalize:
34             cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
35
36         thresh = cm.max() / 2.
37         for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
38             plt.text(j, i, cm[i, j],
39                     horizontalalignment="center",
40                     color="white" if cm[i, j] > thresh else "black")
41
42         plt.tight_layout()
43         plt.ylabel('True label')
44         plt.xlabel('Predicted label')
45

```

```

1 # Creating a simple deep neural network
2 def nn_model(input_dim):
3     model = Sequential()
4     model.add(Dense(256, input_dim=input_dim))
5     model.add(Activation('relu'))
6     model.add(Dropout(0.5))
7
8     model.add(Dense(128))
9     model.add(Activation('relu'))
10    model.add(Dropout(0.5))
11
12    model.add(Dense(64))
13    model.add(Activation('relu'))
14    model.add(Dropout(0.5))
15
16    model.add(Dense(1))
17    model.add(Activation('sigmoid'))
18
19    model.compile(loss='binary_crossentropy', optimizer=Adam(), metrics=['accuracy'])
20    return model

```

```

1 # Concatenating the dataset to increase the dataset
2 x_train_nn = np.concatenate([x_train,x_test])
3 y_train_nn = np.concatenate([y_train,y_test])
4
5 # Normalizing the dataset
6 transformer = Normalizer()
7 X_train_norm = transformer.fit_transform(x_train_nn)
8 X_train_norm = x_train_nn
9
10 # Using 20% data for testing
11 x_test_nn = X_train_norm[int(0.8*len(X_train_norm)):]
12 y_test_nn = y_train_nn[int(0.8*len(X_train_norm)):]
13
14 # Using 80% data for training
15 x_train_nn = X_train_norm[:int(0.8*len(X_train_norm))]
16 y_train_nn = y_train_nn[:int(0.8*len(X_train_norm))]
17
18
19 callbacks = [EarlyStopping(monitor='val_loss', min_delta=0, patience=100, verbose=0, mode='auto'),
20              ModelCheckpoint('output/{val_loss:.4f}.hdf5', monitor='val_loss', verbose=2, save_best_only=True, mode='auto')]

```

```
21
22 # Creating model
23 nn = nn_model(x_train_nn.shape[1])
24
25 # Training. Using 20% of data for validation
26 history = nn.fit(x_train_nn, y_train_nn, validation_split=0.2, epochs=50,
27                 verbose=1, batch_size=8, shuffle=True)
28
```



Train on 128 samples, validate on 32 samples

Epoch 1/50

128/128 [=====] - 1s 5ms/step - loss: 0.6825 - acc: 0.5625 - val\_loss: 0.7178 - val\_acc: 0.3750

Epoch 2/50

128/128 [=====] - 0s 901us/step - loss: 0.6725 - acc: 0.5781 - val\_loss: 0.7493 - val\_acc: 0.3750

Epoch 3/50

128/128 [=====] - 0s 786us/step - loss: 0.6417 - acc: 0.6562 - val\_loss: 0.7913 - val\_acc: 0.3750

Epoch 4/50

128/128 [=====] - 0s 863us/step - loss: 0.6413 - acc: 0.6719 - val\_loss: 0.8197 - val\_acc: 0.3750

Epoch 5/50

128/128 [=====] - 0s 826us/step - loss: 0.6230 - acc: 0.6562 - val\_loss: 0.8762 - val\_acc: 0.3750

Epoch 6/50

128/128 [=====] - 0s 779us/step - loss: 0.6432 - acc: 0.6406 - val\_loss: 0.8738 - val\_acc: 0.3750

Epoch 7/50

128/128 [=====] - 0s 761us/step - loss: 0.6282 - acc: 0.6719 - val\_loss: 0.8906 - val\_acc: 0.3750

Epoch 8/50

128/128 [=====] - 0s 852us/step - loss: 0.6217 - acc: 0.6719 - val\_loss: 0.9242 - val\_acc: 0.3750

Epoch 9/50

128/128 [=====] - 0s 830us/step - loss: 0.6298 - acc: 0.6797 - val\_loss: 0.8764 - val\_acc: 0.3750

Epoch 10/50

128/128 [=====] - 0s 848us/step - loss: 0.6160 - acc: 0.7344 - val\_loss: 0.8894 - val\_acc: 0.3750

Epoch 11/50

128/128 [=====] - 0s 791us/step - loss: 0.5845 - acc: 0.6953 - val\_loss: 0.9018 - val\_acc: 0.3750

Epoch 12/50

128/128 [=====] - 0s 860us/step - loss: 0.6157 - acc: 0.6797 - val\_loss: 0.8825 - val\_acc: 0.3750

Epoch 13/50

128/128 [=====] - 0s 1ms/step - loss: 0.6025 - acc: 0.7031 - val\_loss: 0.8927 - val\_acc: 0.3750

Epoch 14/50

128/128 [=====] - 0s 966us/step - loss: 0.6015 - acc: 0.7188 - val\_loss: 0.8902 - val\_acc: 0.3750

Epoch 15/50

128/128 [=====] - 0s 871us/step - loss: 0.6353 - acc: 0.6641 - val\_loss: 0.8891 - val\_acc: 0.3750

Epoch 16/50

128/128 [=====] - 0s 836us/step - loss: 0.6080 - acc: 0.6797 - val\_loss: 0.8934 - val\_acc: 0.3750

Epoch 17/50

128/128 [=====] - 0s 756us/step - loss: 0.5992 - acc: 0.7031 - val\_loss: 0.8919 - val\_acc: 0.3750

Epoch 18/50

128/128 [=====] - 0s 775us/step - loss: 0.6026 - acc: 0.6875 - val\_loss: 0.9190 - val\_acc: 0.3750

Epoch 19/50

128/128 [=====] - 0s 795us/step - loss: 0.6146 - acc: 0.6641 - val\_loss: 0.8777 - val\_acc: 0.3750

Epoch 20/50

128/128 [=====] - 0s 839us/step - loss: 0.6270 - acc: 0.6875 - val\_loss: 0.8961 - val\_acc: 0.3750

Epoch 21/50

128/128 [=====] - 0s 876us/step - loss: 0.5860 - acc: 0.7266 - val\_loss: 0.8982 - val\_acc: 0.3750



Epoch 22/50

Epoch 22/50  
128/128 [=====] - 0s 871us/step - loss: 0.6375 - acc: 0.7344 - val\_loss: 0.8786 - val\_acc: 0.3750  
Epoch 23/50  
128/128 [=====] - 0s 767us/step - loss: 0.6067 - acc: 0.6953 - val\_loss: 0.8600 - val\_acc: 0.3750  
Epoch 24/50  
128/128 [=====] - 0s 835us/step - loss: 0.6102 - acc: 0.7031 - val\_loss: 0.8919 - val\_acc: 0.3750  
Epoch 25/50  
128/128 [=====] - 0s 909us/step - loss: 0.6159 - acc: 0.6953 - val\_loss: 0.8661 - val\_acc: 0.3750  
Epoch 26/50  
128/128 [=====] - 0s 835us/step - loss: 0.5874 - acc: 0.7031 - val\_loss: 0.8816 - val\_acc: 0.3750  
Epoch 27/50  
128/128 [=====] - 0s 861us/step - loss: 0.6186 - acc: 0.6875 - val\_loss: 0.8506 - val\_acc: 0.3750  
Epoch 28/50  
128/128 [=====] - 0s 797us/step - loss: 0.6329 - acc: 0.6953 - val\_loss: 0.8307 - val\_acc: 0.3750  
Epoch 29/50  
128/128 [=====] - 0s 769us/step - loss: 0.6344 - acc: 0.7188 - val\_loss: 0.8211 - val\_acc: 0.3750  
Epoch 30/50  
128/128 [=====] - 0s 819us/step - loss: 0.5753 - acc: 0.7109 - val\_loss: 0.8482 - val\_acc: 0.3750  
Epoch 31/50  
128/128 [=====] - 0s 848us/step - loss: 0.5959 - acc: 0.6875 - val\_loss: 0.8694 - val\_acc: 0.3750  
Epoch 32/50  
128/128 [=====] - 0s 813us/step - loss: 0.6092 - acc: 0.6797 - val\_loss: 0.9078 - val\_acc: 0.3750  
Epoch 33/50  
128/128 [=====] - 0s 782us/step - loss: 0.5877 - acc: 0.7031 - val\_loss: 0.9068 - val\_acc: 0.3750  
Epoch 34/50  
128/128 [=====] - 0s 886us/step - loss: 0.6054 - acc: 0.6953 - val\_loss: 0.8961 - val\_acc: 0.3750  
Epoch 35/50  
128/128 [=====] - 0s 867us/step - loss: 0.6338 - acc: 0.7031 - val\_loss: 0.8340 - val\_acc: 0.3750  
Epoch 36/50  
128/128 [=====] - 0s 872us/step - loss: 0.6017 - acc: 0.6953 - val\_loss: 0.8586 - val\_acc: 0.3750  
Epoch 37/50  
128/128 [=====] - 0s 948us/step - loss: 0.6450 - acc: 0.6875 - val\_loss: 0.8476 - val\_acc: 0.3750  
Epoch 38/50  
128/128 [=====] - 0s 803us/step - loss: 0.6416 - acc: 0.6875 - val\_loss: 0.8343 - val\_acc: 0.3750  
Epoch 39/50  
128/128 [=====] - 0s 788us/step - loss: 0.6011 - acc: 0.6953 - val\_loss: 0.8502 - val\_acc: 0.3750  
Epoch 40/50  
128/128 [=====] - 0s 828us/step - loss: 0.6001 - acc: 0.7188 - val\_loss: 0.8525 - val\_acc: 0.3750  
Epoch 41/50  
128/128 [=====] - 0s 936us/step - loss: 0.6447 - acc: 0.6562 - val\_loss: 0.8394 - val\_acc: 0.3750  
Epoch 42/50  
128/128 [=====] - 0s 790us/step - loss: 0.5960 - acc: 0.7109 - val\_loss: 0.8486 - val\_acc: 0.3750  
Epoch 43/50  
128/128 [=====] - 0s 804us/step - loss: 0.5998 - acc: 0.6797 - val\_loss: 0.8658 - val\_acc: 0.3750  
Epoch 44/50

```
Epoch 44/50
128/128 [=====] - 0s 781us/step - loss: 0.6151 - acc: 0.6797 - val_loss: 0.8834 - val_acc: 0.3750
Epoch 45/50
128/128 [=====] - 0s 836us/step - loss: 0.5997 - acc: 0.7109 - val_loss: 0.8693 - val_acc: 0.3750
Epoch 46/50
128/128 [=====] - 0s 870us/step - loss: 0.5891 - acc: 0.7188 - val_loss: 0.8750 - val_acc: 0.3750
Epoch 47/50
128/128 [=====] - 0s 787us/step - loss: 0.6259 - acc: 0.7266 - val_loss: 0.8593 - val_acc: 0.3750
Epoch 48/50
128/128 [=====] - 0s 798us/step - loss: 0.5984 - acc: 0.7031 - val_loss: 0.8562 - val_acc: 0.3750
Epoch 49/50
128/128 [=====] - 0s 835us/step - loss: 0.5907 - acc: 0.6953 - val_loss: 0.8776 - val_acc: 0.3750
Epoch 50/50
128/128 [=====] - 0s 786us/step - loss: 0.5838 - acc: 0.7344 - val_loss: 0.8836 - val_acc: 0.3750
```

```
1 # Plot training & validation accuracy values
2 plt.plot(history.history['acc'])
3 plt.plot(history.history['val_acc'])
4 plt.title('Model accuracy')
5 plt.ylabel('Accuracy')
6 plt.xlabel('Epoch')
7 plt.legend(['Train', 'Test'], loc='upper left')
8 plt.show()
9
10 # Plot training & validation loss values
11 plt.plot(history.history['loss'])
12 plt.plot(history.history['val_loss'])
13 plt.title('Model loss')
14 plt.ylabel('Loss')
15 plt.xlabel('Epoch')
16 plt.legend(['Train', 'Test'], loc='upper left')
17 plt.show()
```






```
1 # Calculating test loss and accuracy
2 score = nn.evaluate(x_test, y_test, verbose=2)
3 print('Test loss:', score[0])
4 print('Test accuracy:', score[1])
5
6 # plotting ROC-AUC curve
7 y_pred_nn = nn.predict_proba(x_test)
8 plot_metrics(y_test, y_pred_nn)
```



Test loss: 0.8740641403198243


Test accuracy: 0.45



```
1 # creating classification data
2 y_class = y_pred_nn.copy()
3 y_class[y_class>0.5]=1
4 y_class[y_class<=0.5]=0
5
6 # Confusing matrix
7 confusion_mtx = confusion_matrix(y_test, y_class)
8 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])
```







## ▼ XGBClassifier

```
1 param_grid = {
2     'n_estimators':[10, 20, 30, 40, 50, 60, 70, 80, 90, 100],
3     'max_depth':[100],
4
5 }
6
7 class XGBClassifier_tmp(XGBClassifier):
8     def predict(self, X):
9         return XGBClassifier.predict_proba(self, X)[:, 1]
10
11 clf = GridSearchCV(
12     XGBClassifier_tmp(learning_rate=0.005, subsample=0.8,
13                       colsample_bytree=0.8, n_jobs=20),
14     param_grid=param_grid, n_jobs=5,
15     scoring='roc_auc',
16     cv=StratifiedKFold(7, shuffle=True, random_state=0),
17     verbose=7)
```


```
1 clf.fit(x_train, y_train)
2 clf.best_estimator
```

```
1 clf_best_estimator_  
↳ Fitting 7 folds for each of 10 candidates, totalling 70 fits  
[Parallel(n_jobs=5)]: Using backend LokyBackend with 5 concurrent workers.  
[Parallel(n_jobs=5)]: Done 22 tasks      | elapsed:    5.7s  
[Parallel(n_jobs=5)]: Done 70 out of 70 | elapsed:   11.8s finished  
XGBClassifier(tmp(base_score=0.5, booster='gbtree', colsample_bylevel=1,  
                  colsample_bynode=1, colsample_bytree=0.8, gamma=0,  
                  learning_rate=0.005, max_delta_step=0, max_depth=100,  
                  min_child_weight=1, missing=None, n_estimators=80, n_jobs=20,  
                  nthread=None, objective='binary:logistic', random_state=0,  
                  reg_alpha=0, reg_lambda=1, scale_pos_weight=1, seed=None,  
                  silent=None, subsample=0.8, verbosity=1)
```


```
1 y_pred_xgb = clf.predict(x_test)  
2 y_class = y_pred_xgb.copy()  
3 y_class[y_class>0.5]=1  
4 y_class[y_class<=0.5]=0  
5 print("Accuracy:",metrics.accuracy_score(y_test.squeeze(1), y_class))  
6 print("Precision:",metrics.precision_score(y_test, y_class))  
7 print("Recall:",metrics.recall_score(y_test, y_class))  
8  
9 # confusing matrix  
10 confusion_mtx = confusion_matrix(y_test, y_class)  
11 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])
```

↳

Accuracy: 0.45  
Precision: 0.4528301886792453  
Recall: 0.48



```
1 # plotting ROC-AUC curve
2 y_pred_nn = nn.predict_proba(x_test)
3 plot_metrics(y_test, y_pred_xgb)
```




## ▼ AdaBoostClassifier

```
1 clf = AdaBoostClassifier(n_estimators=120, learning_rate=0.009, random_state=13,  
2                           base_estimator=DecisionTreeClassifier(max_depth=19, min_samples_leaf=40, max_features=2,  
3                                                                    random_state=13))  
4 clf.fit(x_train, y_train.squeeze(1))
```

```
↳ AdaBoostClassifier(algorithm='SAMME.R',  
                      base_estimator=DecisionTreeClassifier(ccp_alpha=0.0,  
                                                             class_weight=None,  
                                                             criterion='gini',  
                                                             max_depth=19,  
                                                             max_features=2,  
                                                             max_leaf_nodes=None,  
                                                             min_impurity_decrease=0.0,  
                                                             min_impurity_split=None,  
                                                             min_samples_leaf=40,  
                                                             min_samples_split=2,  
                                                             min_weight_fraction_leaf=0.0,  
                                                             presort='deprecated',  
                                                             random_state=13,  
                                                             splitter='best'),  
                      learning_rate=0.009, n_estimators=120, random_state=13)
```

```
1 y_pred_ada = clf.predict_proba(x_test)[: , 1]  
2 plot_metrics(y_test, y_pred_ada)
```


```
↳
```




```
1 y_class = y_pred_ada.copy()
2 y_class[y_class>0.5]=1
3 y_class[y_class<=0.5]=0
4 print("Accuracy:",metrics.accuracy_score(y_test.squeeze(1), y_class))
5 print("Precision:",metrics.precision_score(y_test, y_class))
6 print("Recall:",metrics.recall_score(y_test, y_class))
7
8 # confusing matrix
9 confusion_mtx = confusion_matrix(y_test, y_class)
10 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])
```



Accuracy: 0.45  
Precision: 0.4444444444444444  
Recall: 0.4



```
1 # Adding Noise
2 y_pred_ada = add_noise(clf.predict_proba(x_test)[: , 1])
3 plot_metrics(y_test, y_pred_ada)
```



## ▼ GradientBoostingClassifier

```
1 %time
2 gb = GradientBoostingClassifier(learning_rate=0.1, n_estimators=50, subsample=0.8, random_state=13,
3                               min_samples_leaf=1, max_depth=3)
4 gb.fit(x_train_scale, y_train)
```

```
↳ CPU times: user 4 µs, sys: 0 ns, total: 4 µs
Wall time: 8.58 µs
GradientBoostingClassifier(ccp_alpha=0.0, criterion='friedman_mse', init=None,
                           learning_rate=0.1, loss='deviance', max_depth=3,
                           max_features=None, max_leaf_nodes=None,
                           min_impurity_decrease=0.0, min_impurity_split=None,
                           min_samples_leaf=1, min_samples_split=2,
                           min_weight_fraction_leaf=0.0, n_estimators=50,
                           n_iter_no_change=None, presort='deprecated',
                           random_state=13, subsample=0.8, tol=0.0001,
                           validation_fraction=0.1, verbose=0,
                           warm_start=False)
```

```
1 proba_gb = gb.predict_proba(x_test_scale)
2 proba_class = proba_gb.copy()
3 proba_class = proba_class[:, 1]
4 proba_class[proba_class>0.5]=1
5 proba_class[proba_class<=0.5]=0
6
7 print ("Log Loss:", log_loss(y_test, proba_gb))
8 print("Accuracy:",metrics.accuracy_score(y_test.squeeze(1), proba_class))
9 print("Precision:",metrics.precision_score(y_test, proba_class))
10 print("Recall:",metrics.recall_score(y_test, proba_class))
11
12 # confusing matrix
13 confusion_mtx = confusion_matrix(y_test, y_class)
14 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])
```


↳

Log Loss: 1.3733057294268327


Accuracy: 0.52

Precision: 0.52

Recall: 0.52



```
1 plot_metrics(y_test, proba_gb[:,1])
```





## ▼ Support Vector Machine


```
1  clf = svm.SVC(kernel='linear') # Linear Kernel
2  clf.fit(x_train_scale, y_train)
3  y_pred_svm = clf.predict(x_test_scale)
4  print("Accuracy:",metrics.accuracy_score(y_test, y_pred_svm))
5  print("Precision:",metrics.precision_score(y_test, y_pred_svm))
6  print("Recall:",metrics.recall_score(y_test, y_pred_svm))
7
8  # confusing matrix
9  confusion_mtx = confusion_matrix(y_test, y_class)
10 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])
```



Accuracy: 0.46

Precision: 0.46296296296296297

Recall: 0.5



## ▼ KNeighborsClassifier

```
1  from sklearn.neighbors import KNeighborsClassifier
2  clf = KNeighborsClassifier()
3  clf.fit(x_train_scale, y_train)
4  y_pred_knn = clf.predict(x_test_scale)
5  print("Accuracy:",metrics.accuracy_score(y_test, y_pred_knn))
6  print("Precision:",metrics.precision_score(y_test, y_pred_knn))
```


```

7 print("Recall:",metrics.recall_score(y_test, y_pred_knn))
8
9 # confusing matrix
10 confusion_mtx = confusion_matrix(y_test, y_pred_knn)
11 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])

```



Accuracy: 0.49  
Precision: 0.48936170212765956  
Recall: 0.46



## ▼ LogisticRegression


```

1 from sklearn.linear_model import LogisticRegression
2 clf = LogisticRegression()
3 clf.fit(x_train_scale,y_train)
4 y_pred_lr = clf.predict(x_test_scale)
5 print("Accuracy:",metrics.accuracy_score(y_test, y_pred_lr))
6 print("Precision:",metrics.precision_score(y_test, y_pred_lr))
7 print("Recall:",metrics.recall_score(y_test, y_pred_lr))
8
9 # confusing matrix
10 confusion_mtx = confusion_matrix(y_test, y_pred_lr)
11 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])

```



Accuracy: 0.47  
Precision: 0.4716981132075472  
Recall: 0.5




## ▼ DecisionTreeClassifier

```
1 clf = DecisionTreeClassifier()
2 clf.fit(x_train,y_train)
3 y_pred_dt = clf.predict(x_test)
4 print("Accuracy:",metrics.accuracy_score(y_test, y_pred_dt))
5 print("Precision:",metrics.precision_score(y_test, y_pred_dt))
6 print("Recall:",metrics.recall_score(y_test, y_pred_dt))
7
8 # confusing matrix
9 confusion_mtx = confusion_matrix(y_test, y_pred_dt)
10 plot_confusion_matrix(confusion_mtx, ['Signal','Background'])
```



Accuracy: 0.52  
Precision: 0.52  
Recall: 0.52



## ▼ RandomForestClassifier

```
1 from sklearn.ensemble import RandomForestClassifier
2 clf = RandomForestClassifier(n_estimators=1000)
3 clf.fit(x_train, y_train)
4 y_pred_rfc = clf.predict(x_test)
5 print("Accuracy:", metrics.accuracy_score(y_test, y_pred_rfc))
6 print("Precision:", metrics.precision_score(y_test, y_pred_rfc))
7 print("Recall:", metrics.recall_score(y_test, y_pred_rfc))
8
9 # confusing matrix
10 confusion_mtx = confusion_matrix(y_test, y_pred_dt)
11 plot_confusion_matrix(confusion_mtx, ['Signal', 'Background'])
```



Accuracy: 0.52  
Precision: 0.52  
Recall: 0.52

